

Data Distillation for Controlling Specificity in Dialogue Generation

Jiwei Li, Will Monroe and Dan Jurafsky

Computer Science Department

Stanford University, Stanford, CA, USA,

jiweil, wmonroe4, jurafsky@stanford.edu

Abstract

People speak at different levels of specificity in different situations.¹ A conversational agent should have this ability and know when to be specific and when to be general.

We propose an approach that gives a neural network-based conversational agent this ability. Our approach involves alternating between *data distillation* and model training : removing training examples that are closest to the responses most commonly produced by the model trained from the last round and then retrain the model on the remaining dataset. Dialogue generation models trained with different degrees of data distillation manifest different levels of specificity.

We then train a reinforcement learning system for selecting among this pool of generation models, to choose the best level of specificity for a given input. Compared to the original generative model trained without distillation, the proposed system is capable of generating more interesting and higher-quality responses, in addition to appropriately adjusting specificity depending on the context.

Our research constitutes a specific case of a broader approach involving training multiple subsystems from a single dataset distinguished by differences in a specific property one wishes to model. We show that from such a set of subsystems, one can use reinforcement learning to build a system that tailors its output to different input contexts at test time.

1 Introduction

People use different levels of specificity in their language depending on many factors about the context of a conversation: one’s interlocutor, one’s mood, how familiar one is with the topic discussed, how well one understands the other’s utterances, and so forth all influence the decision to respond with generics or specifics. A good dialogue agent should have a similar ability to vary the level of specificity of the responses it generates in an input-dependent way.

When humans speak, we can imagine that each has a series of language models in his mind, each of which is able to generate a sensible response, but which differ in specificity. One picks the appropriate model according to the current situation (whether one understands the input utterance, whether one is interested in the topic, etc.) and generates a dialogue utterance using the selected model. Motivated by this line of thinking, we ask whether a conversational agent could consider a pool of dialogue models that vary in language specificity and pick the best one for producing a response to any given input.

One seemingly straightforward approach would be to split the training data by language specificity and train separate generation models on each split. However, this requires classifying data by text specificity, a problem which poses significant challenges. Language specificity has been historically studied for noun phrases, and a few specificity-indicative features have been identified, such as singular terms, negations, or actual/non-actual moods (Enç, 1991; Lyons, 1995). However, there is no generally agreed criterion for defining the level of specificity of an arbitrary unit of natural language, let alone automatically generating sequences to have different levels of specificity.

In this paper, we propose an iterative *data distillation* approach for addressing this issue.²

¹Depending on their knowledge, interlocutors, mood, etc.

²The model is inspired by the concept of distillation in

The proposed system operates as follows: a neural sequence-to-sequence generation (SEQ2SEQ) model is first trained and used to generate (decode) responses to inputs in a dataset. A list of the most common responses is constructed, and training examples with outputs that are semantically close to these common responses are removed (distilled). This process is then repeated by training another SEQ2SEQ model (from scratch) on the remaining data, decoding using the trained model, collecting generic responses and distilling more data. As the process iterates, responses that are generic are gradually distilled, and the trained models gradually increase in specificity.

At the end of the entire data distillation process, we are presented with a pool of generation models, all of which are able to produce sensible responses to input messages but differ in degree of specificity. This pool of models is analogous to specificity-varying models in a human’s mind. When presented with an input dialogue message, the dialogue system needs to pick one model out of the pool. Which model to choose depends on how well the bot understands the input message, how knowledgeable it is regarding the topic discussed, etc.³ To imbue the agent with this ability, we use reinforcement learning to train a model to pick the an appropriate level of specificity by selecting one of pre-trained generative models from the pool.

Experimental results show that models trained from different rounds of data distillation exhibit a clear spectrum of specificity. Models trained in early rounds of data distillation yield better responses. We also show that the reinforcement learning model is able to choose levels of specificity that are appropriate for a variety of inputs.

Our research constitutes a specific case of a broader approach involving training multiple subsystems from a single dataset distinguished by differences in a specific property one wishes to model (here, specificity), especially when this property is hard to model in a supervised learning setting. We show that from such a set of subsystems, one can use reinforcement learning to build a system that tailors its output to different input contexts at test time.

chemistry, which separates chemical mixtures by gradually increasing the temperature to a point at which one or more compounds in the mixture will vaporize.

³We leave handling other factors that should influence specificity (such as the current mood of the bot and non-linguistic characteristics of the interlocutor) for future work.

2 Related Work

Generic responses in open-domain dialogue

End-to-end dialogue systems (Ritter et al., 2011; Serban et al., 2016c; Vinyals and Le, 2015; Serban et al., 2016d,a; Asghar et al., 2016; Mei et al., 2016), tend to generate highly generic and commonplace responses (Sordoni et al., 2015; Mou et al., 2016). The goal of controlling output specificity is closely related to recent attempts to address this issue. Li et al. (2016a) propose using mutual information as an alternative training objective function in place of maximum likelihood, in which an N-best list generated by $p(t|s)$ is reranked by the backward probability $p(s|t)$.

The aim of this work is more general: instead of attempting to always avoid generic responses, our goal is to provide the system with the flexibility to generate responses at different levels of specificity. Blindly avoiding generating generic responses does not reflect how humans speak: we do say dull, generic things like *I don’t know what you are talking about*, to communicate that we indeed do not understand part of the conversation, or to dismiss something as incorrect or nonsensical. A good dialogue system should have the ability to decide when to say generic things and when not to.

Data manipulation The idea of training with data distillation is inspired by a variety of work in the active learning and subdata selection literature, the key idea of which is to select a subset of a large dataset to train a classifier with minimal performance loss, for when the training dataset is extremely large or training is extremely time-intensive (Wei et al., 2015; Zheng et al., 2014; Prasad et al., 2014; Ghahramani, 2013; Iyer and Bilmes, 2013). The proposed system differs from these subdata selection methods in both goals and implementation: we combine a series of models trained on different subsets of data, with the goal of increasing model performance rather than preserving the model’s performance while reducing the size of the training data.

The system we propose is also related to data manipulation strategies such as boosting (Breiman, 1996b), a type of ensemble method (Dietterich, 2002; Zhou et al., 2002; Krogh et al., 1995) that uses subsets of the original data to produce a series of models and then “boosts” their performance by combining them together; and bagging (Breiman, 1996a), which generates additional data for training using the original dataset to produce multisets

of the same size as the original data, decreasing training variance.

3 Data Distillation

In the section, we describe the proposed data distillation model in detail. We use OpenSubtitles (Tiedemann, 2009) as our training dataset.⁴

3.1 Distilling common responses

We first use the following simple example to illustrate the core idea of our system: consider a model that predicts a multinomial distribution over an output variable (e.g., which fruit to choose). The probability of picking *apple* is 0.3, *orange* 0.25, *blueberry* 0.15, *blackberry* 0.15, and *raspberry* 0.15. Outputs that are generic are usually highly probable, since the high diversity of specific outputs results in each having smaller probability mass. We thus treat *apple* as the most generic fruit, and the various berries as more specific. Maximum likelihood estimation at test time will lead the model to always choose *apple*, since it has the largest probability. Observing that *apple* is the most common output, we will remove all *apples* from the training set and retrain the model, which will pick *orange* this time, since it has the greatest probability after *apples* are removed. We then remove *oranges* and repeat the process. With successive iterations of this distillation process, we will gradually obtain models that produce more specific outputs.

In the context of dialogue response generation, our approach works as follows: for each iteration, we first train a SEQ2SEQ model using attention (Bahdanau et al., 2014; Luong et al., 2015) on the original training set. Next, we use the trained model to decode responses to a number of input examples. We decode only a subset of the training set, 1 million responses in total. One could also use a held-out dataset for decoding, but the source of input messages is fairly unimportant in identifying the most frequent responses. We use greedy decoding (beam search with beam size 1). We then collect the most common responses in a list, denoted by L . A response is considered generic if its frequency of occurrence exceeds a threshold,

⁴OpenSubtitles is a large, noisy, open-domain dataset of lines from movie scripts. The noise in the dataset is largely due to the lack of speaker labels for lines of the subtitles. Following Vinyals et al. (2015), we train our models to predict the current line given the preceding ones, assuming that each line constitutes a full speaker turn and that consecutive turns belong to the same conversation. Both assumptions are occasionally untrue but yield reasonable results.

Input: training data D
Output: sequence of trained models M

```

 $M \leftarrow \emptyset$ 
for  $i \leftarrow 1$  to  $N = 8$  do
  train a SEQ2SEQ model  $m$  on  $D$  until convergence
   $M \leftarrow M + m$ 
  decode subset of input messages in  $D$  using model  $m$ 
  collect top frequent decoded responses  $L$ 
  for all instances  $e \in D$  do
    compute relevance score  $R(e)$  using Eq. 1
  end for
   $D^\neg \leftarrow$  top examples by  $R(e)$ 
  distill  $D^\neg$ :  $D \leftarrow D - D^\neg$ 
end for
return  $M$ 

```

Algorithm 1: A brief summary of the proposed data distillation algorithm.

which is empirically set to 100 in this work. We then compare each response in the training data to each highly frequent response from the list L and assign a relevance score $R(e)$ to each training example e based on the cosine similarity between e and the sequence most similar to it in L :

$$R(e) = \max_{e' \in L} \cos(e, e') \quad (1)$$

We use the encoder part of the trained encoder-decoder model to map these sequences to vector representations, which are used to compute the cosine similarity. In this way, sentences that are semantically similar to frequent responses are assigned high relevance scores.⁵ We then remove (distill) examples from the training data with the highest relevance scores⁶ and retrain a new SEQ2SEQ model on the data that remains. An outline of the distillation algorithm is shown in Algorithm 1.

3.2 Choosing a specificity model

The data distillation process produces a pool of SEQ2SEQ models, each trained on the dataset remaining after a different data distillation round. When presented with an input message at test time, the system has to decide which generation model from the pool to use to decode a response to the input. We repeat the data distillation process 8 times, which means we have 8 models in the pool to choose from.⁷ The system should have the ability to choose different models in response to properties

⁵Other options include skip-thought vectors (Kiros et al., 2015) and bag-of-word representations. We find using the trained encoder works decently well.

⁶The amount to remove is empirically set to 8–10%.

⁷It requires two Tesla K40 GPUs to fit the 8 models in memory.

of different inputs. For example, a good dialogue system should give concrete responses when asked things that it is sure about, but generic ones when the input message is difficult to understand.

We use reinforcement learning to train a model to make this choice. Given an input message X from a held-out dataset, we parameterize the action of choosing the generative model with index i from the pool $G = \{g_i\}$ of SEQ2SEQ models trained with data distillation as a policy network $\pi(g_i|X)$, which produces a distribution over $|G|$ classes. To compute the distribution, we first map the input X to a vector representation h_X using an LSTM and then map h_X to a policy distribution over different $g_i \in G$ using a softmax function:

$$\pi(g = g_i|X) = \frac{\exp(h_X^T \cdot h_{g_i})}{\sum_{j=1}^{|G|} \exp(h_X^T \cdot h_{g_j})} \quad (2)$$

where h_{g_i} is an output vector for each model g_i that is randomly initialized and then trained. Given an action, namely a choice of a generative model g_i , we start decoding given the input message X using that model. Decoding generates an output response y , which yields a reward $R(y)$ evaluating response quality according to some metric. The reward signal $R(y)$ is used to train the policy network.

We use the REINFORCE algorithm (Williams, 1992), a kind of policy gradient method, to find the optimal policy by maximizing the expected reward $E_{\pi(g_i|X)}[R(y)]$. The expectation is approximated by sampling from π and the gradient is computed using the likelihood ratio (Aleksandrov et al., 1968):

$$\nabla E(\theta) \approx [R(y) - b] \nabla \log \pi(g_i|X) \quad (3)$$

where b denotes a baseline value.⁸

Adversarial evaluation for reward calculation

One remaining question is how to assign a reward R to a generated response y given the input X , which boils down to the fundamental question of how to evaluate the general quality of a generated response. Dialogue quality is traditionally evaluated (Sordoni et al., 2015, e.g.) using word-overlap metrics such as BLEU and METEOR scores used for machine translation, which have recently been found to correlate poorly with human evaluations (Liu et al., 2016). Recent work has begun using

⁸The baseline value is estimated using another neural model. We refer the readers to Ranzato et al. (2015) and Zaremba and Sutskever (2015) for more details.

more flexible and reliable evaluation metrics; automatic prediction of human ratings (Lowe et al., 2016) is one such metric, but this approach requires a large amount of human labeling effort to train a prediction model.

We employ adversarial evaluation (Li et al., 2016c; Anjuli and Oriol, 2016) for reward calculation. The idea of adversarial evaluation, first proposed by Bowman et al. (2015), is to train a discriminator (or evaluator) function to labels dialogues as machine-generated (negative) or human-generated (positive), a binary classification task. For our system, we use positive examples taken directly from training dialogues, while negative examples are decoded using generative models from different rounds of data distillation. To be specific, for each input message, we randomly sample a SEQ2SEQ model from the pool to decode a response to the input and use the response as a negative example. The evaluator is a **hierarchical neural model** (Serban et al., 2016b): dialogue utterances (i.e., source messages and responses) are first mapped to vector representations using an LSTM. Another LSTM is applied to the sequence of utterance representations to produce **a dialogue representation, which is then fed to a binary classifier**.

Given a pre-trained evaluator D , an input source X and a machine generated target y decoded by the chosen generative model, the **reward R used to update the policy π is the probability that the evaluator D assigns to labeling y as a human-generated response**. The policy update influences the choice of generative model for decoding the current input X . We refer readers to (Li et al., 2016c) for more details about the adversarial evaluation.

3.3 Stochastic Greedy Sampling

Language specificity also relates to language diversity. Utterances with lower levels of diversity are usually generic **because generic responses are usually generic in the same way**. Modeling diversity also provides an indirect way to handle the issue of specificity.

Moreover, there is a degree of **randomness** in human language production: in the real world, if we ask a person the same question twice, even with the same environment and surroundings, it is unlikely that the person will give the same answer both times. **Sampling from the distribution** not only better mimic the way humans generate tokens, but also provides a way to handle the issue of language specificity.

One simple solution is to sample directly from the distribution $p(y|x)$ in all cases. However, we observe that sampling leads to incoherent, ungrammatical, or even irrelevant responses. We expect there to be a sweet spot on the spectrum of randomness, between full sampling on one end and greedy or beam search on the other.⁹

We propose a straightforward algorithm called **Stochastic Greedy Sampling**, in which instead of sampling from the full distribution over all candidate tokens, the model only samples from the few (e.g., 5) words with the highest probability. The model provides with both the flexibility of incorporating randomness and the rigidity of adhering to a pre-trained generation model at the same time.

Again, we use Adversarial Evaluation for comparing purposes. We report *AdverSuc* and *machine-vs-random* proposed by Anjuli and Oriol (2016). *machine-vs-random* denotes the the accuracy of distinguishing between machine-generated responses and randomly sampled responses using a machine evaluator, trained in a way similar to the evaluator in *AdverSuc*. Table 1 presents results for *AdverSuc* and *machine-vs-random* results for *greedy decoding*, *pure sampling* and the proposed *stochastic greedy model*. As can be seen, *sampling* all the time obtains the best score for *AdverSuc*, but also extremely low score for *machine-vs-random* accuracy, which indicates the inferiority of the always sampling strategy. The proposed *stochastic greedy* model perform better than always taking greedy actions as in *greedy*. This indicates that properly combining greedy search and sampling will potentially lead to better results.

Model	AdverSuc	machine-vs-random
greedy	0.042	0.935
pure sampling	0.384	0.642
stochastic greedy	0.058	0.933

Table 1: Adversarial evaluation results for different greedy vs. sampling decoding strategies.

4 Experimental Results

In this section, we present the results of experiments.

⁹Since greedy decoding has been shown to generate higher-quality responses than beam search in dialogue response generation (Li et al., 2016a), we focus on greedy decoding. However, all algorithms can be easily adapted to use beam-search decoding.

iter	data size	ppl	oracle-ppl	div-1	div-2
1	45.2M	33.2	33.2	0.65%	1.57%
2	40.6M	33.3	32.3	0.92%	2.81%
3	35.7M	33.7	31.6	1.18%	3.22%
4	32.7M	34.3	31.2	1.44%	3.60%
5	30.1M	35.0	30.8	1.87%	3.94%
6	27.9M	35.5	30.7	2.21%	4.32%
7	25.5M	36.7	30.5	2.72%	4.65%
8	22.8M	37.2	30.3	3.10%	5.01%

Table 2: Training set size (examples) after data distillation in each iteration and perplexity (ppl) and n -gram diversity scores (dis- n) of the trained generative models on the development set.

4.1 Comparing generative models from different iterations

It is interesting to first compare the generative models and the remaining training data from each of the 8 rounds of data distillation. We use *Iter+N* to denote the generation model trained on the dataset after N repetitions of data distillation.

Perplexity and diversity The size of the training dataset after each round of data distillation and the perplexity of the corresponding trained models on the full development set is shown in the first two columns of Table 2. Perplexity increases for models trained with more data distillation (as expected, since distillation removes opportunities for the model to learn to produce the most common outputs).

However, we expect models trained with distillation to complement the model trained on the entire dataset by better modeling more specific outputs. To quantify the potential of the pool of generation models to complement each other when used in different contexts, we also report **oracle perplexity** (“oracle-ppl”) as a function of the number of iterations K : for each example, we identify the generation model (out of *Iter1* through *IterK*) that assigns the highest probability to the true output. Oracle perplexity is the perplexity computed using these maximal probabilities, instead of the probabilities assigned by any one model. This is equivalent to the perplexity of a model with an RL policy network that chooses perfectly every time. We expect to find that oracle perplexity on the development set decreases when adding the models trained in the first few rounds of data distillation, after which it levels off. This confirms that there are benefits to be had from choosing smartly among the different models.

Count	Response	Count	Response
Iter1		Iter2	
145575	i don't know what you are talking about .	54227	i 'm not in the mood .
84435	i 'm not going to let you go .	29559	i 'm sorry about the way i acted .
36032	i 'm sorry i didn't mean to offend you .	22987	you're not in the mood .
23890	i 'm not so sure .	21392	i 'm gonna take a look at the new york times .
19405	i don't know what to say .	20380	i 'll be there in a minute .
16888	i 'm not going to let you go !	14736	i 'm gonna take a look at this .
16048	that's a good idea .	13753	i 'll get the money .
12782	i don't know what to do .	13013	i 'm gonna take a shower .
11840	i 'm not going to be able to do that .	11746	i 'm in the middle of a war .
11604	i 'm sorry i can't help you .	10130	you're not getting any sleep .
11254	i 'm sure you're right .	9996	i 'm gonna take a look at the other side .
9474	you don't know what you are saying .	9644	i 'm sorry about the way you did .
9471	i 'm not going to tell you .	9169	i 've been doing a lot of things .
8905	i 'm not sure i can do it .	7837	you're a dead man .
7905	i have no idea .	5320	i was just getting a little tired of it .
Iter3		Iter4	
41139	i 'm not an idiot .	30378	i 'm not from around here .
34738	i 'm not an expert on this .	26705	i 'm not from the future .
20252	i 'm sorry but i 'm not an expert on this .	9923	i was just talking to my wife .
16275	i 've got some bad news for you .	9012	i 'm not doing this .
16081	i 'll get you a new suit .	8573	you're a goddamn liar .
13007	i 'm not an idiot !	7424	i 'll be on the way .
11254	i 'm gonna make a big deal out of this .	6919	i 'm sorry ma'am .
6532	i 'm just an ordinary man .	5546	i 'm going back to the hotel .
5724	i 'm not an expert on the police .	4569	i 'll be on my way .
5604	i 'm not an expert on the subject .	4555	i 'm not staying here .
5168	i 'm not your enemy !	4416	you're a goddamn genius .
4963	i 'm not an expert on the law .	4184	i 'm a little tired .
4454	i 'm gonna need some more help with this .	4183	i 'm gonna take a look at this .
4342	i was just about to get my hands on the wall .	4103	he's a bit of a jerk .
3969	i can't believe you're still alive .	3819	he's a bit of a pain in the ass .

Table 3: Most frequent responses generated using greedy search at the end of 1–4 rounds of data distillation. “Count” indicates the number of occurrences of a response in 1 million decoded outputs.

Table 2 also shows a measure of the diversity of generated responses, namely, the number of distinct unigrams (“div-1”) and bigrams (“div-2”) in generated responses as a fraction of the total generated tokens, as described in (Li et al., 2016a). As can be seen, as the data distillation process proceeds and more generic responses are distilled, the system generates increasingly diverse responses.

Distilled responses The highest-frequency responses from different rounds of data distillation are shown in Table 3. Top responses are more generic for models trained in earlier iterations. In iteration 1, the top responses are broadly generic statements of uncertainty (“*I don’t know*”, “*I am not sure*”) or agreement (“*i think you are right*” or “*that’s a good idea*”), but the meanings of frequent responses start diverging as the distillation algorithm proceeds. The number of the occurrences of the top frequent responses from different iterations also validates this point, with the number gradually decreasing.

Table 4 presents sampled outputs from the generation models trained after different rounds of

the distillation. Responses from *Iter1* are usually generic but safe, mostly *i don’t know what to do/what you are talking about* and *that’s a good idea*. As the amount of distilled data increases, the corresponding model generates increasingly concrete responses but has a greater risk of **outputting confusing or irrelevant responses**.

4.2 Choosing the correct model for decoding

Next, we present results from the proposed reinforcement learning model and analyze how it decides which model to pick from the pool.

The distribution over different models used to decode input messages in the development set is shown in Figure 1. As can be seen, the RL model chooses to decode using the model trained on the entire dataset (i.e., *Iter1*) for 16 percent of all inputs. The models trained after 2, 3 and 4 rounds are responsible for decoding responses to approximately half of the inputs.

Human evaluation For human evaluation, we follow protocols defined in Li et al. (2016b), employing crowdsourced judges to evaluate a random

<i>Input:</i> hear it ? Iter1: i 'm sorry i didn 't hear you . <i>Iter5:</i> that 's the only one who could ever be .	<i>RL model choice:</i> Iter1 <i>Iter2:</i> it 's a little early in the morning for me . <i>Iter7:</i> what 's the meaning of this ?
<i>Input:</i> it feels like i must have been asleep for weeks . Iter1: i 'm sorry i can 't help you . <i>Iter5:</i> and then i felt sorry about it .	<i>RL model choice:</i> Iter1 <i>Iter2:</i> it 's been so long . <i>Iter7:</i> i 'm sorry i woke you up
<i>Input:</i> we lived in railroad apartments inside red brick . Iter1: oh, really ? <i>Iter5:</i> we got to get back to the hotel and get a job .	<i>RL model choice:</i> Iter1 <i>Iter2:</i> we got a lot of stuff in the trunk . <i>Iter7:</i> i lived in a hotel .
<i>Input:</i> you know you ought to find yourself a girl . Iter1: i 'm not going to do that . <i>Iter5:</i> i 'm gonna make you some of your own kind of girl .	<i>RL model choice:</i> Iter1 <i>Iter2:</i> you know i 'm a real looker . <i>Iter7:</i> you 're in a big house !
<i>Input:</i> the man who does that is not afraid to die . Iter1: i am not afraid of him . <i>Iter5:</i> the man is a man of god who is afraid of god	<i>RL model choice:</i> Iter1 <i>Iter2:</i> but he 's not afraid of the truth . <i>Iter7:</i> it is my duty to protect the father .
<i>Input:</i> the cellar he was working on something in the cellar . <i>Iter1:</i> i was in the cellar . <i>Iter5:</i> there 's no way he could have gotten into this room .	<i>RL model choice:</i> Iter2 Iter2: i 'm going to the cellar . <i>Iter7:</i> we had the whole room .
<i>Input:</i> this little lady seems to be alone in the world . <i>Iter1:</i> i 'm not going to be able to see her . <i>Iter5:</i> you know i can tell when she is a little bit but she 's not very good	<i>RL model choice:</i> Iter2 Iter2: i 'm sure she 'll have a good time with the rest of the crew . <i>Iter7:</i> but she 's not coming here .
<i>Input:</i> i didn 't even see you go back there . <i>Iter1:</i> i was just going to see you . <i>Iter5:</i> i got to see your father .	<i>RL model choice:</i> Iter2 Iter2: i didn 't see the man . <i>Iter7:</i> i went back to sleep .
<i>Input:</i> i was kinda hoping that you were my dad . <i>Iter1:</i> i was just wondering if you could help me out . <i>Iter5:</i> he 's been a long time since he got here .	<i>RL model choice:</i> Iter2 Iter2: you 're the only one who 's ever been in my family . <i>Iter7:</i> so what are you doing with my mom ?
<i>Input:</i> excuse me but how did we pay for these suits ? <i>Iter1:</i> i don 't know what you are talking about . <i>Iter5:</i> i can 't afford this .	<i>RL model choice:</i> Iter2 Iter2: i don 't know if i can get a refund on this . <i>Iter7:</i> i 'll pay them 10 000 .
<i>Input:</i> maybe his eyes were too big for his stomach. <i>Iter1:</i> he was a big boy . <i>Iter5:</i> i think he was a good kid .	<i>RL model choice:</i> Iter2 Iter2: he 's got a big mouth . <i>Iter7:</i> he was just trying to kill her .
<i>Input:</i> like the death zone or the zone of no return . <i>Iter1:</i> i don 't know what to do . <i>Iter5:</i> the man is a man of the universe .	<i>RL model choice:</i> Iter2 Iter2: the zone is the only place where the people are . <i>Iter7:</i> but it 's a life zone .
<i>Input:</i> home sir . <i>Iter1:</i> that 's great ! Iter5: the rest of the crew are also on their way home .	<i>RL model choice:</i> Iter5 <i>Iter3:</i> the rest of the world is here . <i>Iter7:</i> how 's your wife ? .
<i>Input:</i> i 'm sorry to detain you for so long . <i>Iter1:</i> i 'm sorry i can 't . Iter5: i 'm not leaving until tomorrow .	<i>RL model choice:</i> Iter5 <i>Iter2:</i> i 'm sorry about the delay . <i>Iter7:</i> i 'm glad that 's your name and all .
<i>Input:</i> don 't fuckin talk to me that way man . <i>Iter1:</i> i 'm not gonna talk to you . <i>Iter5:</i> i 'm just sayin that i 'm an asshole and that ain 't the way i talk !	<i>RL model choice:</i> Iter7 <i>Iter2:</i> i 'm talkin to you motherfucker . Iter7: you 're in no position to talk like that .
<i>Input:</i> can 't you see how they 're exploiting you ? <i>Iter1:</i> i 'm not a human being . <i>Iter5:</i> you 're not an idiot anymore .	<i>RL model choice:</i> Iter7 <i>Iter2:</i> you 're not a good guy . Iter7: well it 's not that simple .
<i>Input:</i> do you have any idea what caused the explosion ? <i>Iter1:</i> i have no idea . <i>Iter5:</i> i 've never heard of any of these .	<i>RL model choice:</i> Iter7 <i>Iter2:</i> i think it 's the same as a gas pipe . Iter7: the explosion was caused by a meteor disturbance .

Table 4: Sample of outputs generated by models from different rounds (1,2,5,7) of data distillation, along with the model that the proposed reinforcement learning (RL) algorithm chooses to decode the response. Responses chosen by the RL model are in **Bold**.

sample of 200 items. We present labelers with an input message and the generated outputs from three models, *Iter1*, *Iter2*, and *RL*, and ask them to rank the three outputs by quality. Note that the outputs from the *RL* model can be the same as those from *Iter1* or *Iter2* if the *RL* model chooses that partic-

ular model (*Iter1* or *Iter2*) for decoding. In these cases, a tie is automatically recorded. Figure 2 shows the proportions of the outputs ranked first by the human labelers. As can be seen, the reinforcement learning model performs best 60 percent of the time, followed by *Iter2*, which wins 37 percent




Figure 1: Distribution of the iteration used by the RL model to decode responses (dev set).




Figure 2: The proportions of outputs from different models ranked first in the human evaluation. Note that the sum is larger than 100 percent due to ties.

pairs	win	lose	tie
RL vs. Iter2	51	29	20
RL vs. Iter1	64	28	8
Iter2 vs. Iter1	62	38	-

Table 5: Pairwise human judgments between the reinforcement learning model (*RL*) and the first two distillation models (*Iter1* and *Iter2*).

of the time. Table 5 shows pairwise human judgments between the three models extracted from the three-instance ranking. It is interesting to see that *Iter2* generally outperforms *Iter1*, winning on 62 percent of the examples. This is consistent with the fact that the *RL* model tends to prefer *Iter2* more often.

Adversarial evaluation Table 6 reports adversarial success and *machine-vs-random* accuracy described in Li et al. (2016c). Adversarial success (*AdverSuc*) refers to the percentage of machine-generated responses that are able to fool a trained evaluator model into believe that they are generated by humans; *machine-vs-random* accuracy denotes the accuracy of a trained evaluator model (a different evaluator from the one used in adversarial success) at distinguishing between machine-generated responses and human utterances randomly sampled without regard for the input. Superior models should obtain higher values of both adversarial suc-

model	<i>AdverSuc</i>	<i>machine-vs-random</i>
Iter1 (standard)	0.058	0.933
random	0.056	0.940
RL	0.088	0.944

Table 6: Adversarial success and *machine-vs-random* accuracy for *Iter1* which always generating response using the model trained on the full set, *random* which randomly samples a model for generation, and the proposed model.

cess and *machine-vs-random* accuracy. We refer readers to Li et al. (2016c) for more details. We observe that the *RL* model performs better than always using the model trained on the full dataset (*Iter1*) or choosing a distillation model at random (as one would expect, since the *RL* model is trained to optimize adversarial success).

Analyzing results Table 4 shows example choices made of the *RL* model in response to different inputs. When input messages are **vague and hard to reply to**, the *RL* model usually picks *Iter1*, which in turn outputs safe responses like “that’s great” or “i don’t know what you are talking about”. The *RL* model has a tendency to pick models from the latter stages of distillation training if all of the generation models from the different iterations of distillation **are able to output meaningful responses**, since models from the later stages output produce more diverse and interesting outputs. We also observe a high correlation between the number of **unknown words** in the source sentence and the choice to use *Iter1*.

5 Conclusion

In this paper, we investigate the language specificity issue in dialogue generation. We propose a **data distillation method**, which trains a series of generation models that exhibit different levels of specificity and uses a reinforcement learning model to choose the model best suited for decoding depending on the dialogue context.

The success of the proposed system confirms the importance of **data processing** in training a successful open-domain dialogue system. We anticipate that strategies resembling the one we propose can be used more generally for controlling properties of dialogue generation other than specificity, by training several models on different subsets of a single dataset that differ in the desired property, and choosing among these to produce outputs that tailor the quality of interest to the situation at hand.

References

- V. M. Aleksandrov, V. I. Sysoyev, and V. V. Shemeneva. 1968. Stochastic optimization. *Engineering Cybernetics* 5:11–16.
- Kannan Anjuli and Vinyals Oriol. 2016. Adversarial evaluation of dialogue models .
- Nabiha Asghar, Pasca Poupart, Jiang Xin, and Hang Li. 2016. Online sequence-to-sequence reinforcement learning for open-domain conversational agents. *arXiv preprint arXiv:1612.03929* .
- Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. 2014. Neural machine translation by jointly learning to align and translate. *arXiv preprint arXiv:1409.0473* .
- Samuel R Bowman, Luke Vilnis, Oriol Vinyals, Andrew M Dai, Rafal Jozefowicz, and Samy Bengio. 2015. Generating sentences from a continuous space. *arXiv preprint arXiv:1511.06349* .
- Leo Breiman. 1996a. Bagging predictors. *Machine learning* 24(2):123–140.
- Leo Breiman. 1996b. Bias, variance, and arcing classifiers .
- Thomas G Dietterich. 2002. Ensemble learning. *The handbook of brain theory and neural networks* 2:110–125.
- Mürvet Enç. 1991. The semantics of specificity. *Linguistic inquiry* pages 1–25.
- Zoubin Ghahramani. 2013. Scaling the indian buffet process via submodular maximization. *arXiv preprint arXiv:1304.3285* .
- Rishabh K Iyer and Jeff A Bilmes. 2013. Submodular optimization with submodular cover and submodular knapsack constraints. In *Advances in Neural Information Processing Systems*. pages 2436–2444.
- Ryan Kiros, Yukun Zhu, Ruslan R Salakhutdinov, Richard Zemel, Raquel Urtasun, Antonio Torralba, and Sanja Fidler. 2015. Skip-thought vectors. In *Advances in neural information processing systems*. pages 3294–3302.
- Anders Krogh, Jesper Vedelsby, et al. 1995. Neural network ensembles, cross validation, and active learning. *Advances in neural information processing systems* 7:231–238.
- Jiwei Li, Michel Galley, Chris Brockett, Jianfeng Gao, and Bill Dolan. 2016a. A diversity-promoting objective function for neural conversation models. In *Proc. of NAACL-HLT*.
- Jiwei Li, Will Monroe, Alan Ritter, and Dan Jurafsky. 2016b. Deep reinforcement learning for dialogue generation. *arXiv preprint arXiv:1606.01541* .
- Jiwei Li, Will Monroe, Tianlin Shi, and Dan Jurafsky. 2016c. Adversarial reinforcement learning for neural dialogue generation.
- Chia-Wei Liu, Ryan Lowe, Iulian V Serban, Michael Noseworthy, Laurent Charlin, and Joelle Pineau. 2016. How not to evaluate your dialogue system: An empirical study of unsupervised evaluation metrics for dialogue response generation. *arXiv preprint arXiv:1603.08023* .
- Ryan Lowe, Michael Noseworthy, Iulian Serban, Nicolas Angelard-Gontier, Yoshua Bengio, and Joelle Pineau. 2016. Towards an automatic turing test: Learning to evaluate dialogue responses. *arXiv preprint arXiv:1605.05414* .
- Minh-Thang Luong, Hieu Pham, and Christopher D Manning. 2015. Effective approaches to attention-based neural machine translation. *arXiv preprint arXiv:1508.04025* .
- John Lyons. 1995. *Linguistic semantics: An introduction*. Cambridge University Press.
- Hongyuan Mei, Mohit Bansal, and Matthew R Walter. 2016. Coherent dialogue with attention-based language models. *arXiv preprint arXiv:1611.06997* .
- Lili Mou, Yiping Song, Rui Yan, Ge Li, Lu Zhang, and Zhi Jin. 2016. Sequence to backward and forward sequences: A content-introducing approach to generative short-text conversation. *arXiv preprint arXiv:1607.00970* .
- Adarsh Prasad, Stefanie Jegelka, and Dhruv Batra. 2014. Submodular meets structured: Finding diverse subsets in exponentially-large structured item sets. In *Advances in Neural Information Processing Systems*. pages 2645–2653.
- Marc’Aurelio Ranzato, Sumit Chopra, Michael Auli, and Wojciech Zaremba. 2015. Sequence level training with recurrent neural networks. *arXiv preprint arXiv:1511.06732* .
- Alan Ritter, Colin Cherry, and William B Dolan. 2011. Data-driven response generation in social media. In *Proceedings of the conference on empirical methods in natural language processing*. Association for Computational Linguistics, pages 583–593.
- Iulian V Serban, Il Ororbia, G Alexander, Joelle Pineau, and Aaron Courville. 2016a. Multi-modal variational encoder-decoders. *arXiv preprint arXiv:1612.00377* .
- Iulian V Serban, Alessandro Sordani, Yoshua Bengio, Aaron Courville, and Joelle Pineau. 2016b. Building end-to-end dialogue systems using generative hierarchical neural network models. In *Proceedings of AAAI*.
- Iulian Vlad Serban, Ryan Lowe, Laurent Charlin, and Joelle Pineau. 2016c. Generative deep neural networks for dialogue: A short review .

- Iulian Vlad Serban, Alessandro Sordoni, Ryan Lowe, Laurent Charlin, Joelle Pineau, Aaron Courville, and Yoshua Bengio. 2016d. A hierarchical latent variable encoder-decoder model for generating dialogues. *arXiv preprint arXiv:1605.06069*.
- Alessandro Sordoni, Michel Galley, Michael Auli, Chris Brockett, Yangfeng Ji, Meg Mitchell, Jian-Yun Nie, Jianfeng Gao, and Bill Dolan. 2015. A neural network approach to context-sensitive generation of conversational responses. In *Proceedings of NAACL-HLT*.
- Jörg Tiedemann. 2009. News from OPUS – A collection of multilingual parallel corpora with tools and interfaces. In *Recent Advances in Natural Language Processing*, volume 5, pages 237–248.
- Oriol Vinyals and Quoc Le. 2015. A neural conversational model. In *Proceedings of ICML Deep Learning Workshop*.
- Kai Wei, Rishabh Iyer, and Jeff Bilmes. 2015. Submodularity in data subset selection and active learning. In *Proceedings of the 32nd International Conference on Machine Learning, Lille, Fran.* pages 6–11.
- Ronald J Williams. 1992. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine learning* 8(3-4):229–256.
- Wojciech Zaremba and Ilya Sutskever. 2015. Reinforcement learning neural Turing machines. *arXiv preprint arXiv:1505.00521*.
- Jingjing Zheng, Zhuolin Jiang, Rama Chellappa, and Jonathon P Phillips. 2014. Submodular attribute selection for action recognition in video. In *Advances in Neural Information Processing Systems*. pages 1341–1349.
- Zhi-Hua Zhou, Jianxin Wu, and Wei Tang. 2002. Ensembling neural networks: many could be better than all. *Artificial intelligence* 137(1):239–263.